

Applied Regression and Time Series Analysis

Homework 1: OLS Estimation

W271-4

Lei Yang

Subhashini R

Ron Cordell

Question 1:

Load the birthweight dataset. Note that the actual data is provided in a data table named "data". Use the following procedures to load the data:

- Step 1: put the provided R Workspace birthweight w271.RData in the directory of your choice.
- Step 2: Load the dataset using this command: `load("birthweight.Rdata")`

```
# Load the data from the birthweight_w271.Rdata file  
load("birthweight_w271.Rdata")
```

```
library(car)
```

```
ls()
```

```
'data '  
'desc '
```

Question 2:

Examine the basic structure of the data set using `desc`, `str`, and `summary` to examine all of the variables in the data set. How many variables and observations in the data? These commands will be useful:

- `desc`
- `str(data)`
- `summary(data)`

```
desc
```

	variable	label
1	faminc	1988 family income, \$1000s
2	cigtax	cig. tax in home state, 1988
3	cigprice	cig. price in home state, 1988
4	bwght	birth weight, ounces
5	fatheduc	father's yrs of educ
6	motheduc	mother's yrs of educ
7	parity	birth order of child
8	male	=1 if male child
9	white	=1 if white
10	cigs	cigs smoked per day while preg
11	lbwght	log of bwght
12	bwghtlbs	birth weight, pounds
13	packs	packs smoked per day while preg
14	lfaminc	log(faminc)

```
str(data)
```

```
'data.frame':  1388 obs. of  14 variables:
 $ faminc  : num  13.5 7.5 0.5 15.5 27.5 7.5 65 27.5 27.5 37.5 ...
 $ cigtax  : num  16.5 16.5 16.5 16.5 16.5 16.5 16.5 16.5 16.5 16.5 ...
 $ cigprice: num  122 122 122 122 122 ...
 $ bwght   : num  109 133 129 126 134 118 140 86 121 129 ...
 $ fatheduc: int   12 6 NA 12 14 12 16 12 12 16 ...
 $ motheduc: int   12 12 12 12 12 14 14 14 17 18 ...
 $ parity  : int    1 2 2 2 2 6 2 2 2 2 ...
 $ male    : int    1 1 0 1 1 1 0 0 0 0 ...
 $ white   : int    1 0 0 0 1 0 1 0 1 1 ...
 $ cigs     : int    0 0 0 0 0 0 0 0 0 0 ...
 $ lbwght  : num   4.69 4.89 4.86 4.84 4.9 ...
 $ bwghtlbs: num   6.81 8.31 8.06 7.88 8.38 ...
 $ packs    : num    0 0 0 0 0 0 0 0 0 0 ...
 $ lfaminc  : num   2.603 2.015 -0.693 2.741 3.314 ...
- attr(*, "datalabel")= chr ""
- attr(*, "time.stamp")= chr "25 Jun 2011 23:03"
- attr(*, "formats")= chr  "%9.0g" "%9.0g" "%9.0g" "%8.0g" ...
- attr(*, "types")= int   254 254 254 252 251 251 251 251 251 ...
- attr(*, "val.labels")= chr  "" "" "" "" ...
- attr(*, "var.labels")= chr  "1988 family income, $1000s" "cig. tax in home
state, 1988" "cig. price in home state, 1988" "birth weight, ounces" ...
- attr(*, "version")= int 10
```

summary(data)

faminc	cigtax	cigprice	bwght
Min. : 0.50	Min. : 2.00	Min. :103.8	Min. : 0.0
1st Qu.:14.50	1st Qu.:15.00	1st Qu.:122.8	1st Qu.:106.0
Median :27.50	Median :20.00	Median :130.8	Median :119.0
Mean :29.03	Mean :19.55	Mean :130.6	Mean :117.9
3rd Qu.:37.50	3rd Qu.:26.00	3rd Qu.:137.0	3rd Qu.:132.0
Max. :65.00	Max. :38.00	Max. :152.5	Max. :271.0

fatheduc	motheduc	parity	male
Min. : 1.00	Min. : 2.00	Min. :1.000	Min. :0.0000
1st Qu.:12.00	1st Qu.:12.00	1st Qu.:1.000	1st Qu.:0.0000
Median :12.00	Median :12.00	Median :1.000	Median :1.0000
Mean :13.19	Mean :12.94	Mean :1.633	Mean :0.5209
3rd Qu.:16.00	3rd Qu.:14.00	3rd Qu.:2.000	3rd Qu.:1.0000
Max. :18.00	Max. :18.00	Max. :6.000	Max. :1.0000
NA's :196	NA's :1		

white	cigs	lbwght	bwghtlbs
Min. :0.0000	Min. : 0.000	Min. :0.000	Min. : 0.000
1st Qu.:1.0000	1st Qu.: 0.000	1st Qu.:4.663	1st Qu.: 6.625
Median :1.0000	Median : 0.000	Median :4.779	Median : 7.438
Mean :0.7846	Mean : 2.087	Mean :4.726	Mean : 7.366
3rd Qu.:1.0000	3rd Qu.: 0.000	3rd Qu.:4.883	3rd Qu.: 8.250
Max. :1.0000	Max. :50.000	Max. :5.602	Max. :16.938

packs	lfaminc
Min. :0.0000	Min. : -0.6931
1st Qu.:0.0000	1st Qu.: 2.6741
Median :0.0000	Median : 3.3142
Mean :0.1044	Mean : 3.0713
3rd Qu.:0.0000	3rd Qu.: 3.6243
Max. :2.5000	Max. : 4.1744

There are 1388 observations of 14 variables

Question 3:

As we mentioned in the live session, it is important to start with a question (or a hypothesis) when conducting regression modeling. In this exercise, we are in the question: "Do mothers who smoke have babies with lower birth weight?" The dependent variable of interest is `bwght`, representing birthweight in ounces. Examine this variable using both tabulated summary and graphs. Specifically,

1. Summarize the variable `bwght`: `summary(data$bwght)`
2. You may also use the quantile function: `quantile(data$bwght)`. List the following quantiles: 1%, 5%, 10%, 25%, 50%, 75%, 90%, 95%, 99%
3. Plot the histogram of `bwght` and comment on the shape of its distribution. Try different bin sizes and comment how it affects the shape of the histogram. Remember to label the graph clearly. You will also need a title for the graph.
4. This is a more open-ended question: Have you noticed anything "strange" with the `bwght` variable and the shape of histogram this variable? If so, please elaborate on your observations and investigate any issues you have identified.

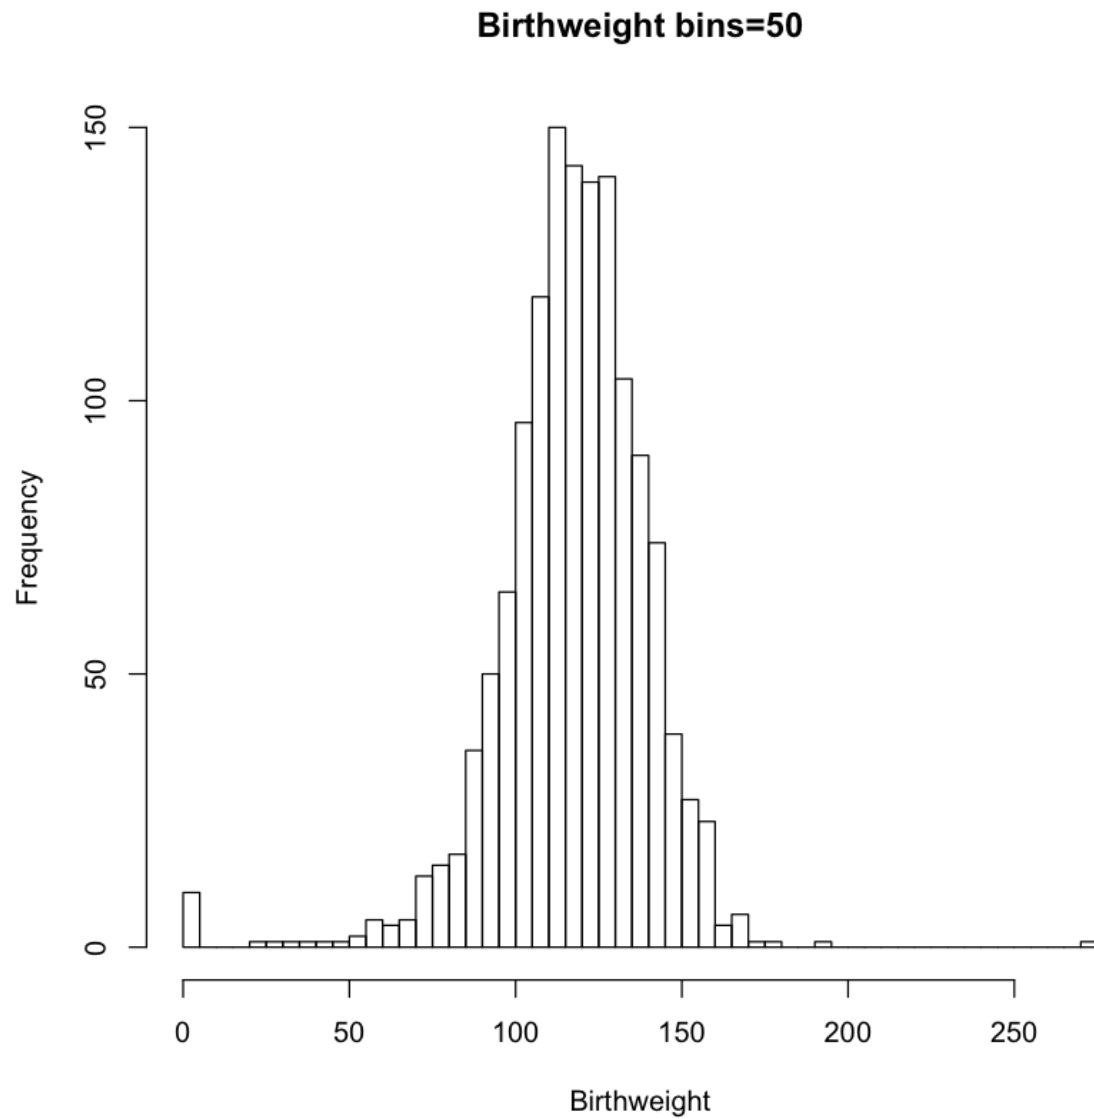
```
summary(data$bwght)
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   0.0   106.0   119.0   117.9   132.0   271.0
```

```
quantile(data$bwght, probs=c(1,5,10,25,50,75,90,99, NA)/100)
```

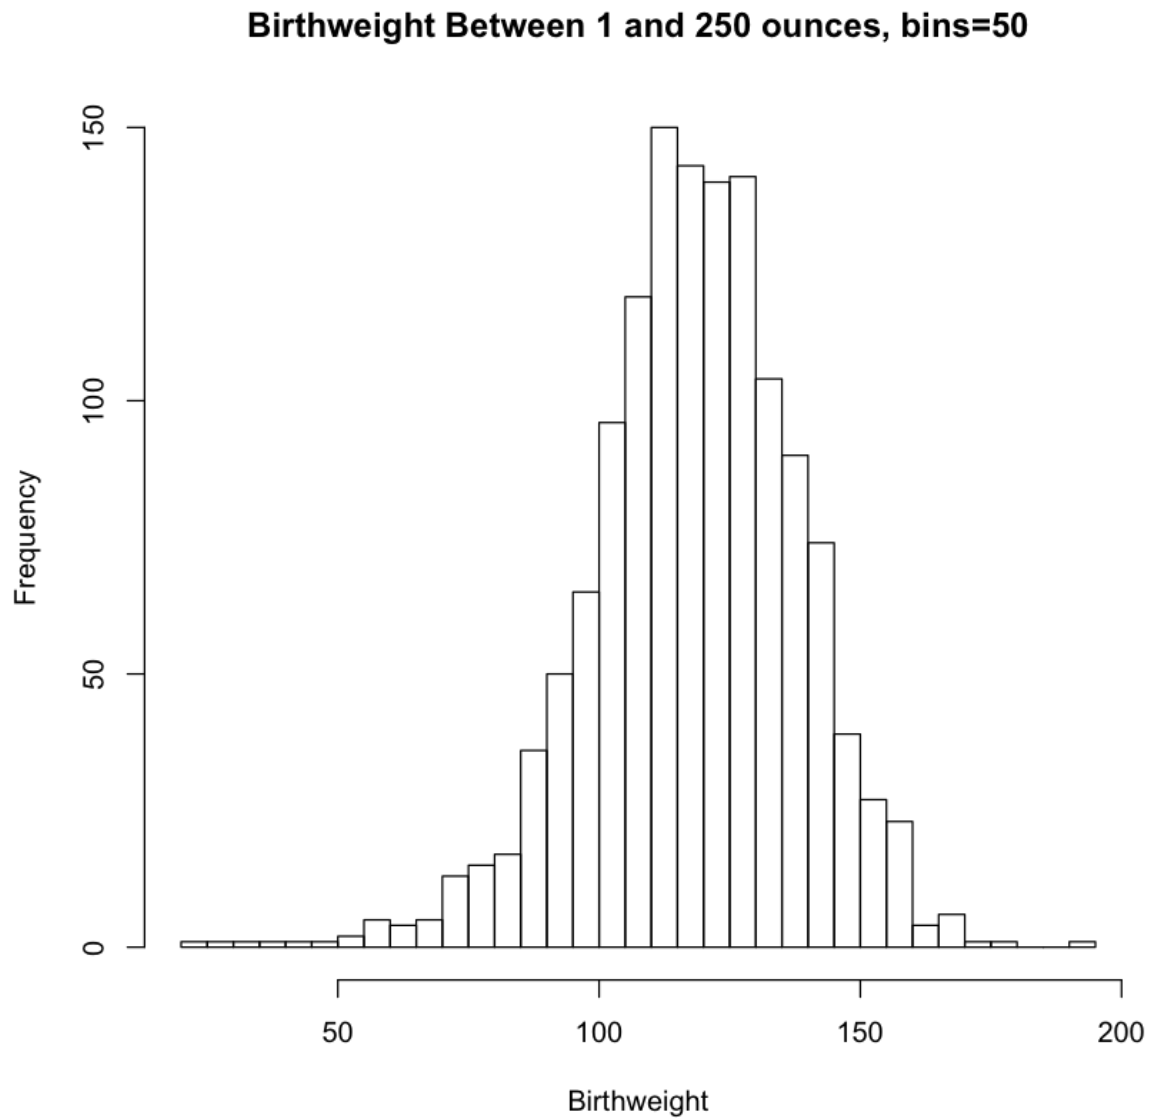
	quantile	count
1	1%	42.35
2	5%	83
3	10%	93
4	25%	106
5	50%	119
6	75%	132
7	90%	143
8	99%	160.13
9	9	NA

```
hist(data$bwght, main='Birthweight bins=50', xlab='Birthweight',  
breaks=50)
```



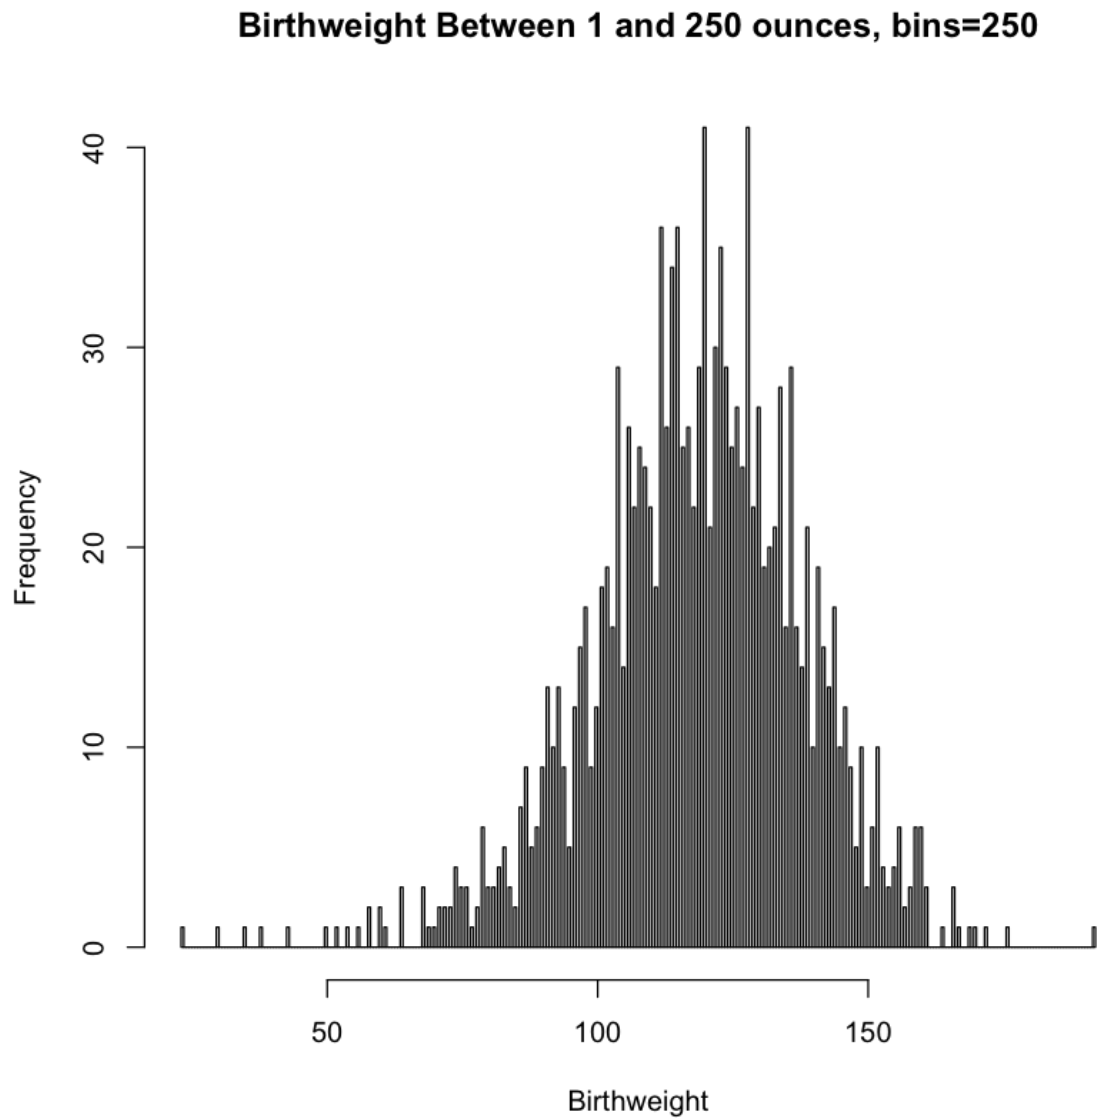
Setting the binning to break=50 reveals finer structure in the data. One obvious point is that there is an outlier data point beyond the 250 mark, which skews the visual appearance of the data. There are also birthweights of 0 which should be treated as NA.

```
hist(data$bwght[data$bwght<250 & data$bwght>0], main='Birthweight  
Between 1 and 250 ounces, bins=50', xlab='Birthweight', breaks=50)
```



Removing the outliers by filtering out those data points greater than 250 shows that the distribution is slightly longer left-tailed but otherwise normal-like.

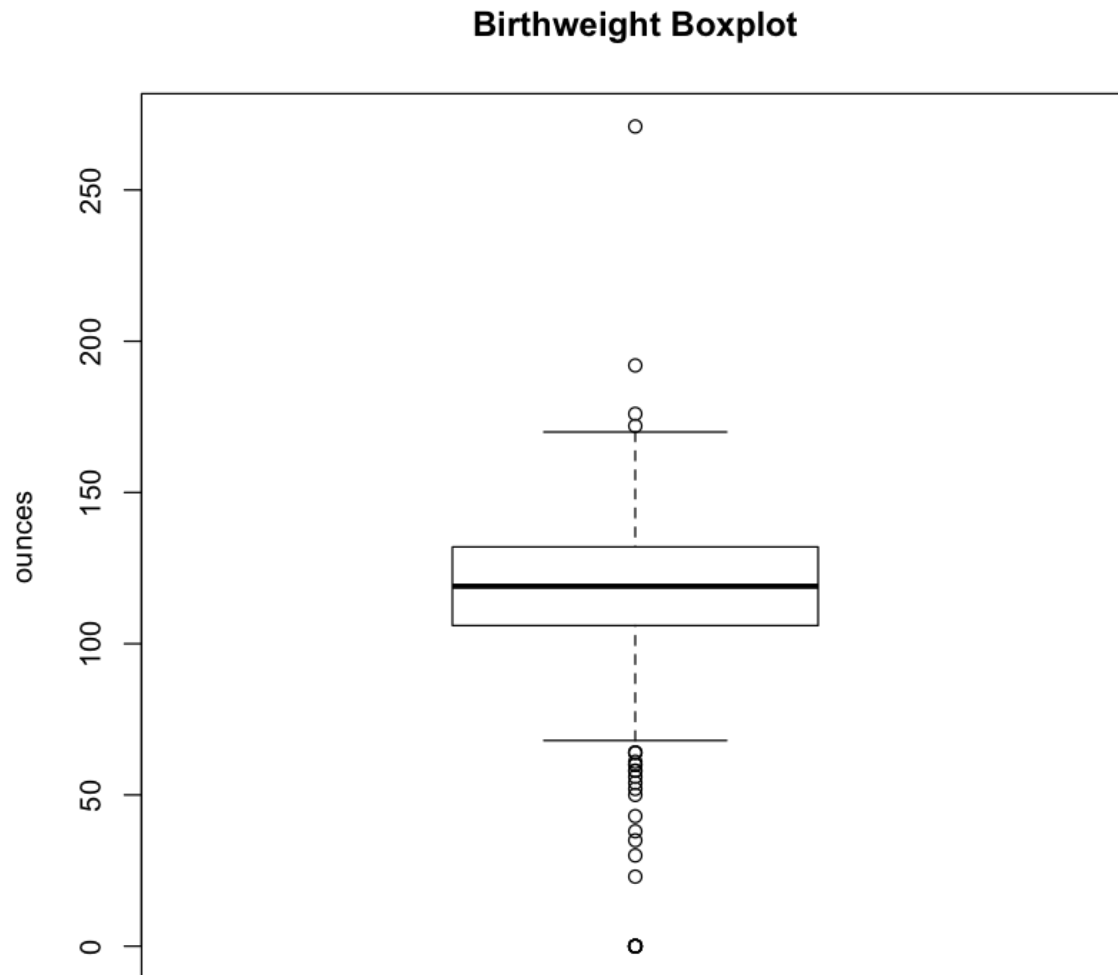
```
hist(data$bwght[data$bwght<250 & data$bwght>0], main='Birthweight  
Between 1 and 250 ounces, bins=250', xlab='Birthweight', breaks=250)
```



As we increase the number of bins we see more variance indicated in the graph.

Boxplot of Birthweight

```
boxplot(data$bwght, data=data, main = "Birthweight Boxplot",  
ylab=("ounces"))
```



The boxplot corresponds the histogram in that it shows that the distribution tends towards the mean with a narrow spread. There is also the longer spread of values on the bottom than on the top, similar to the long left tail of the histogram.

Question 4:

Examine the variable `cigs`, which represents number of cigarettes smoked each day by the mother while pregnant. Conduct the same analysis as in question 3.

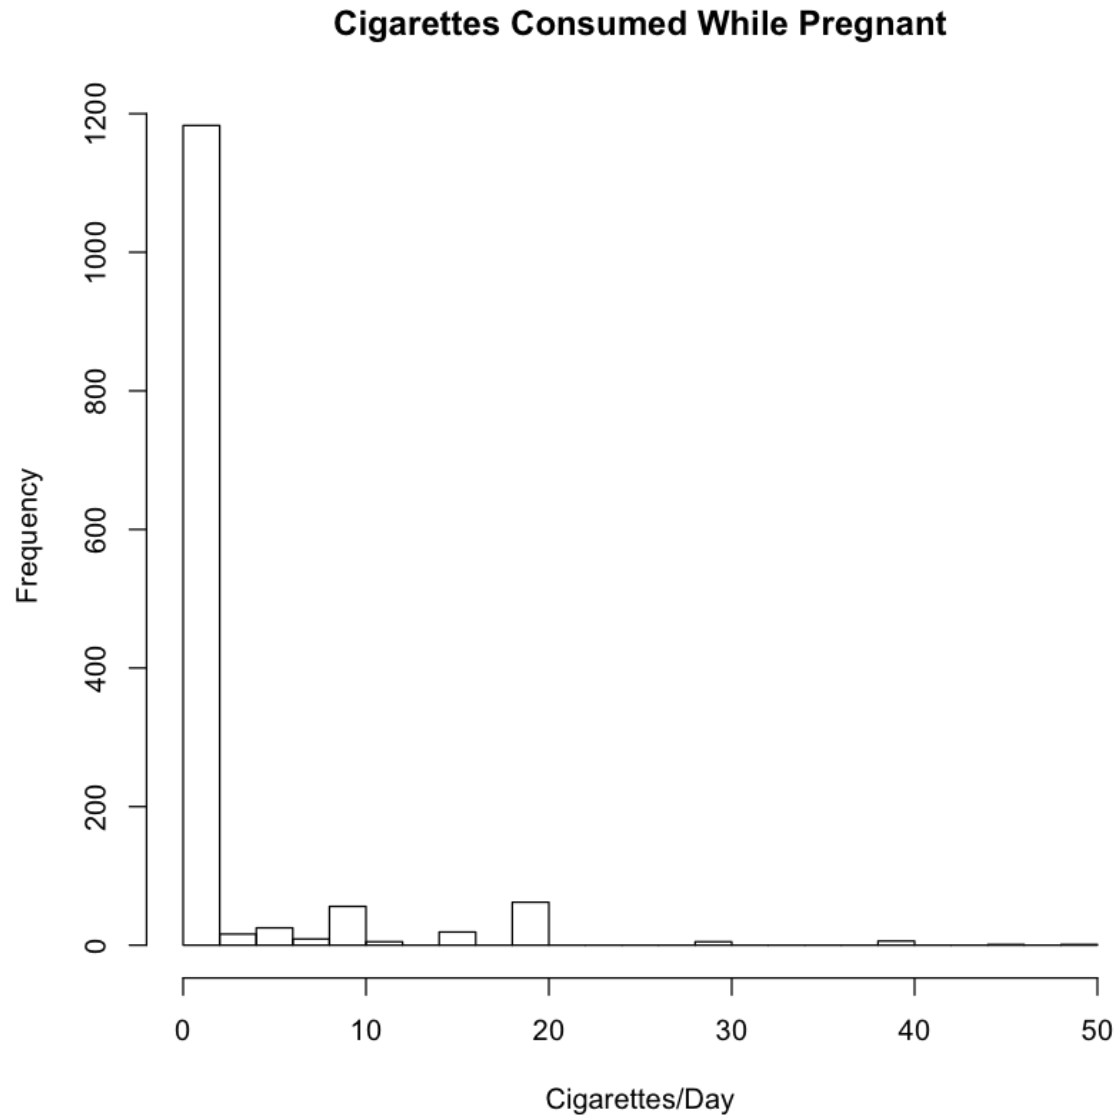
```
summary(data$cigs)
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.000    0.000    0.000    2.087    0.000   50.000
```

```
quantile(data$cigs, probs=c(1,5,10,25,50,75,90,99, NA)/100)
```

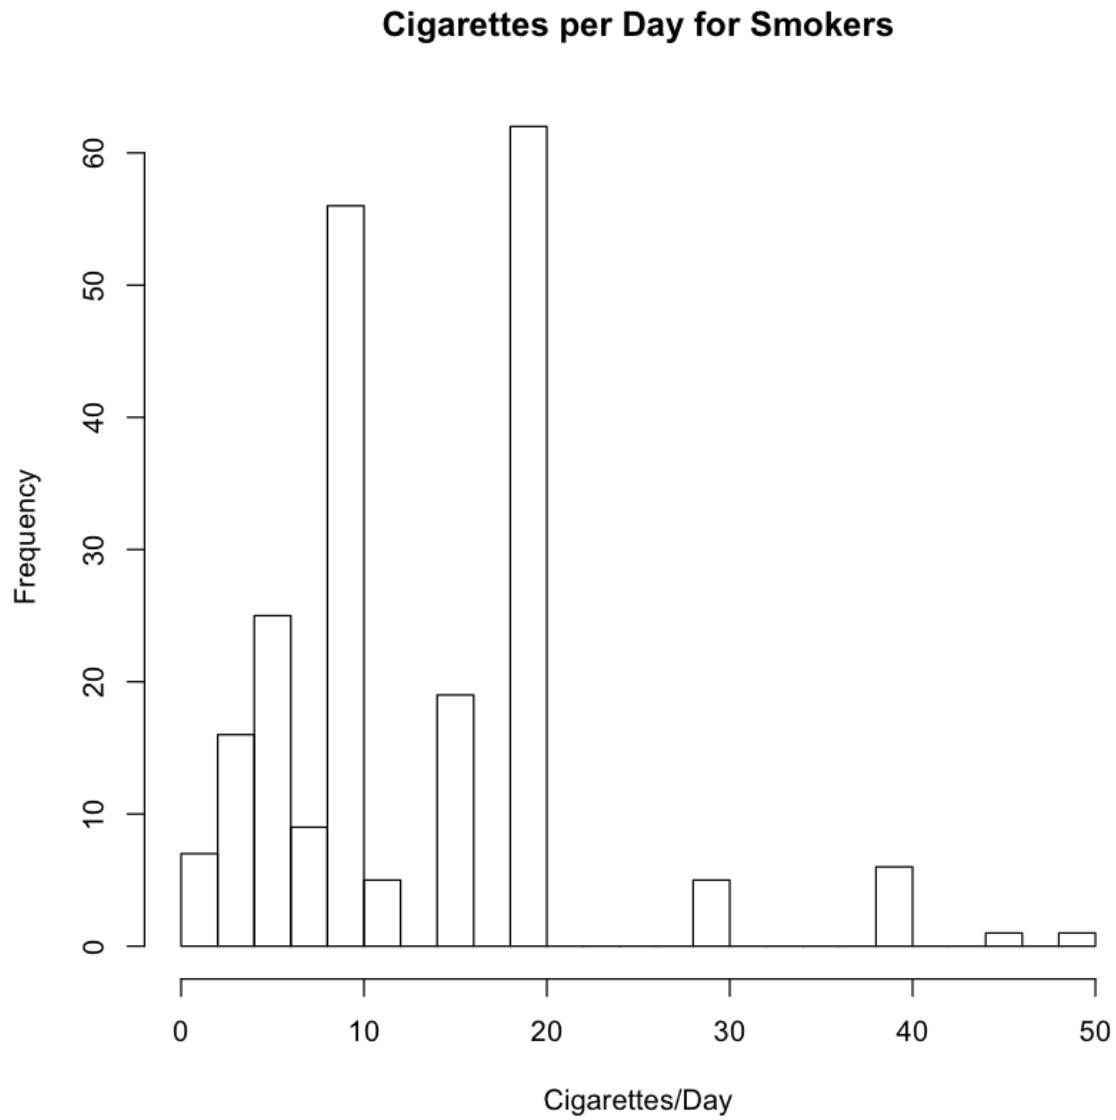
	quantile	count
1	1%	0
2	5%	0
3	10%	0
4	25%	0
5	50%	0
6	75%	0
7	90%	10
8	99%	20
9	9	NA

```
hist(data$cigs, main='Cigarettes Consumed While Pregnant',  
xlab='Cigarettes/Day', breaks=25)
```




The histogram for cigarettes per day shows a very large spike for all the 0 values in the data set. Filtering out the zero values gives a better sense for the relative frequencies of the non-zero values, shown below.

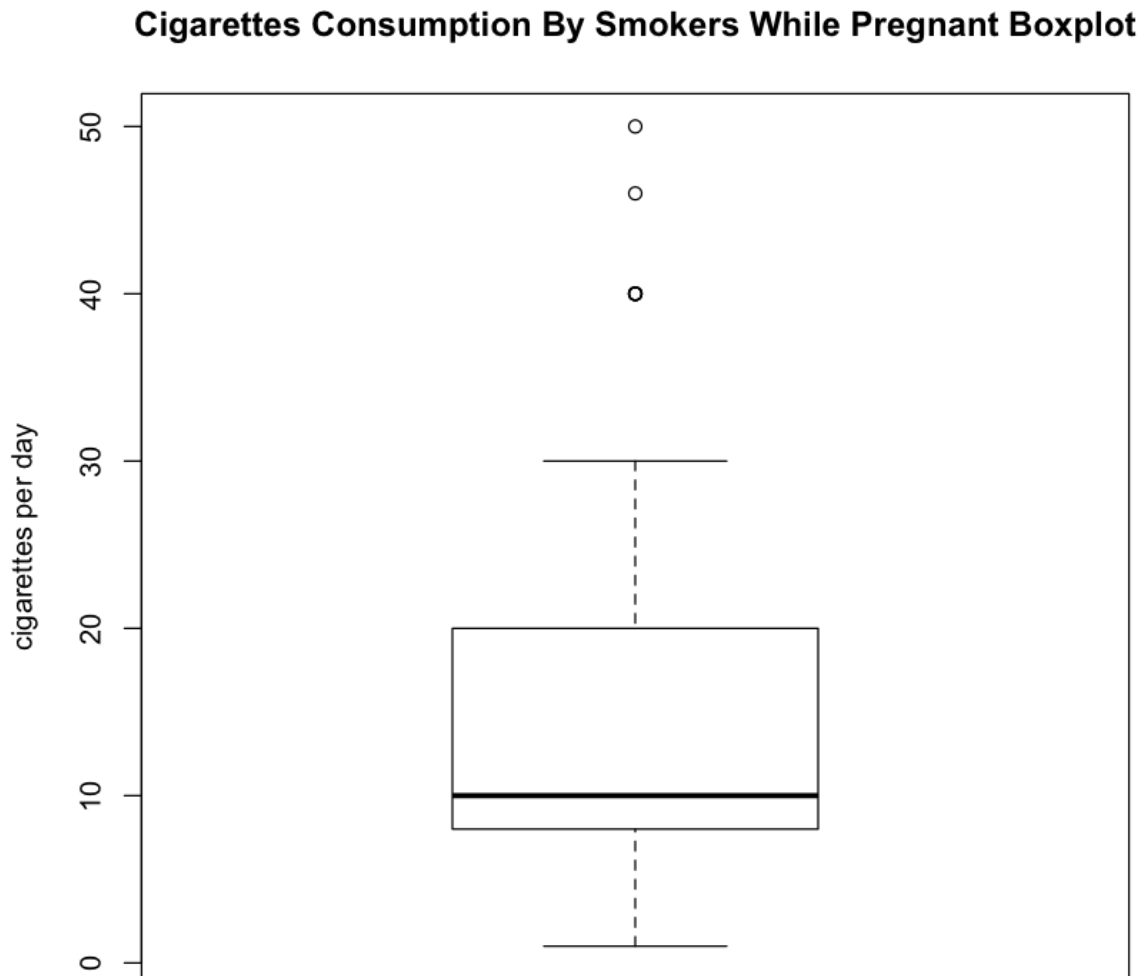
```
hist(data$cigs[data$cigs>0], main='Cigarettes per Day for Smokers',  
xlab='Cigarettes/Day', breaks=25)
```



```
hist(data$cigs[data$cigs>0], main='Cigarettes per Day for Smokers',  
xlab='Cigarettes/Day', breaks=50)
```



```
boxplot(data$cigs[data$cigs>0], data=data, main="Cigarettes  
Consumption By Smokers While Pregnant Boxplot",ylab="cigarettes per  
day")
```

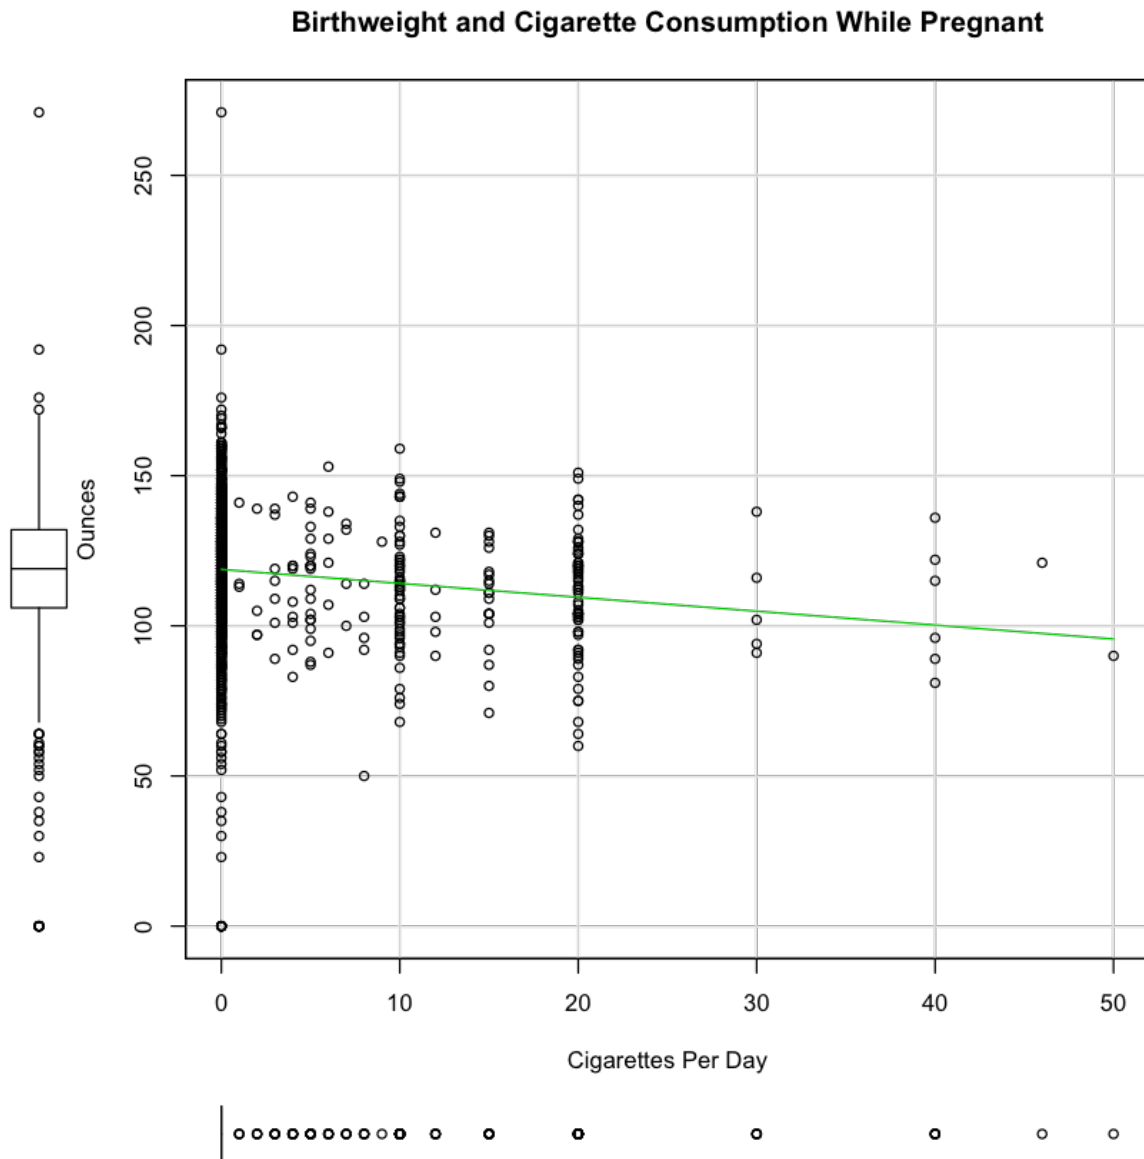


The cigarettes per day boxplot also indicates a wide spread of data with a number of outliers in the upper consumption range. This boxplot filters out the zero cigarettes per day values.

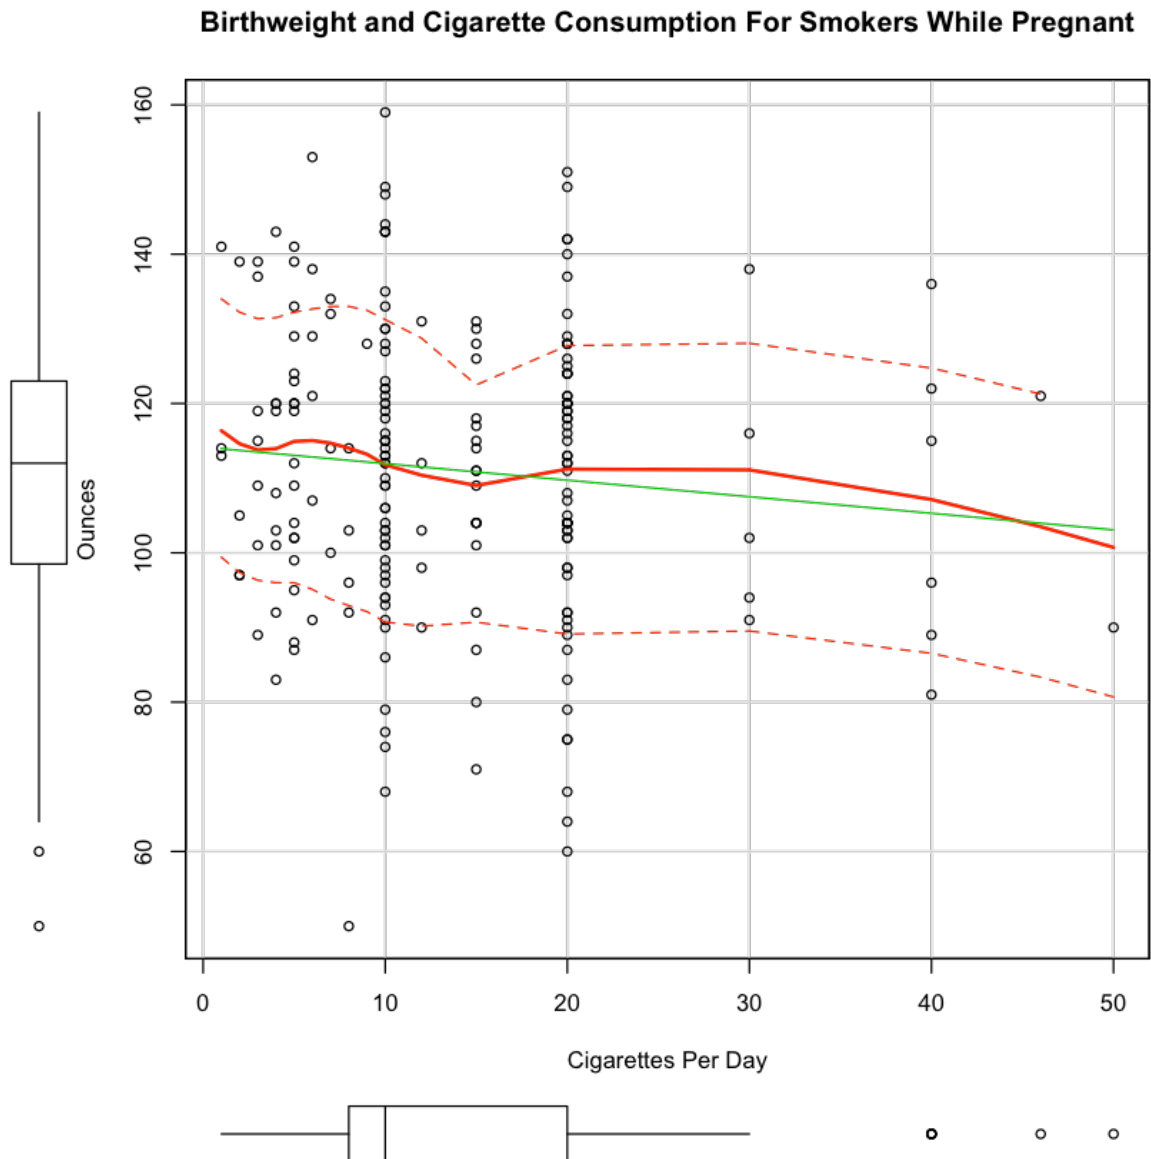
Question 5:

Generate a scatterplot of bwght against cigs. Based on the appearance of this plot, how much of the variation in bwght do you think can be explained by cigs?

```
scatterplot(data$cigs, data$bwght,  
main="Birthweight and Cigarette Consumption While Pregnant",xlab =  
"Cigarettes Per Day", ylab="Ounces")
```



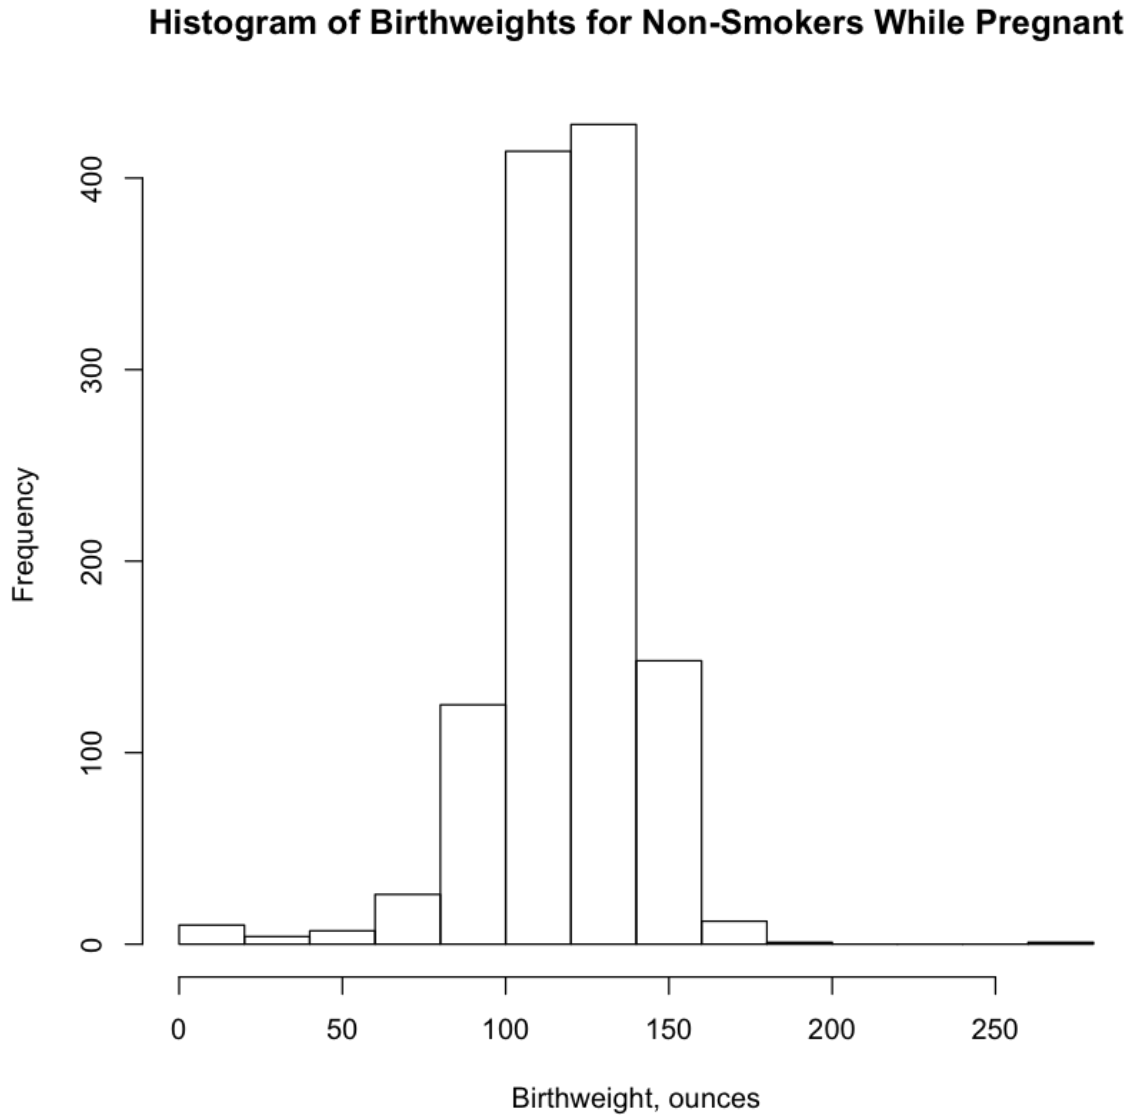
```
scatterplot(data$cigs[data$cigs>0], data$bwght[data$cigs>0],
main="Birthweight and Cigarette Consumption For Smokers While Pregnant",
xlab = "Cigarettes Per Day", ylab="Ounces")
```




There appears to be less variation in birthweight with cigarette consumption but the data is also skewed such that there are many more samples for zero cigarette consumption by two orders of

magnitude. If we replot and filter out the non-smokers there still is a slight negative correlation between cigarette consumption while pregnant and birthweight.

```
hist(data$bwght[data$cigs==0], main="Histogram of Birthweights for  
Non-Smokers While Pregnant", xlab="Birthweight, ounces")
```




```
hist(data$bwght[data$cigs>0], main="Histogram of Birthweights for  
Smokers While Pregnant",  
      xlab="Birthweight, ounces")
```



The two histograms show how the birthweight samples are distributed for smokers and non-smokers. There is a slightly lower mean for the smokers, but the variance is higher.


Question 6:

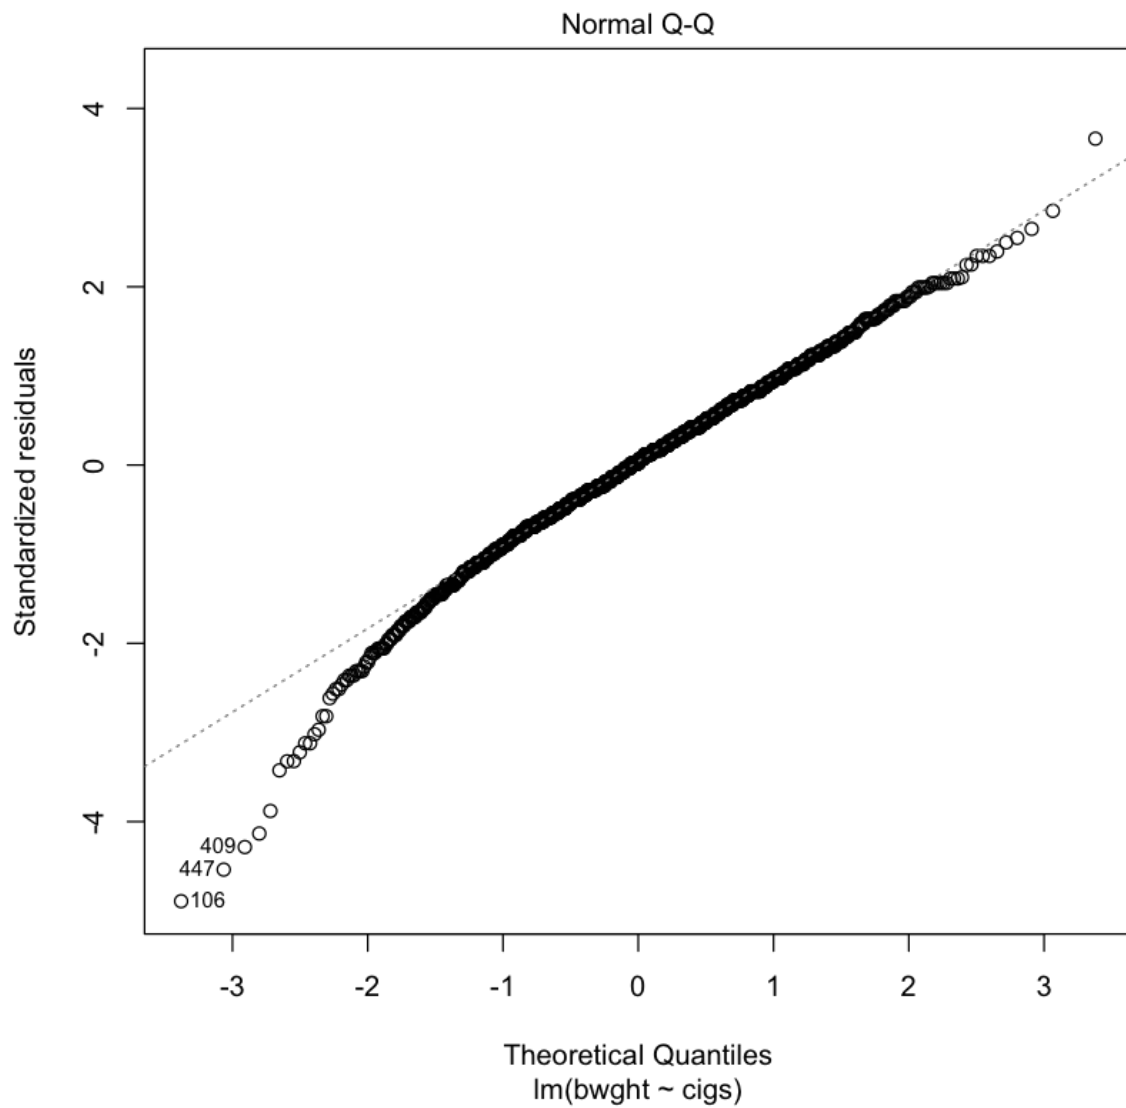
Estimate the simple linear regression of bwght on cigs. What coefficient estimates and the standard errors associated with the coefficient estimates do you get? Interpret the results. Note that you may have to "take care of" any potential data issues before building a regression model.

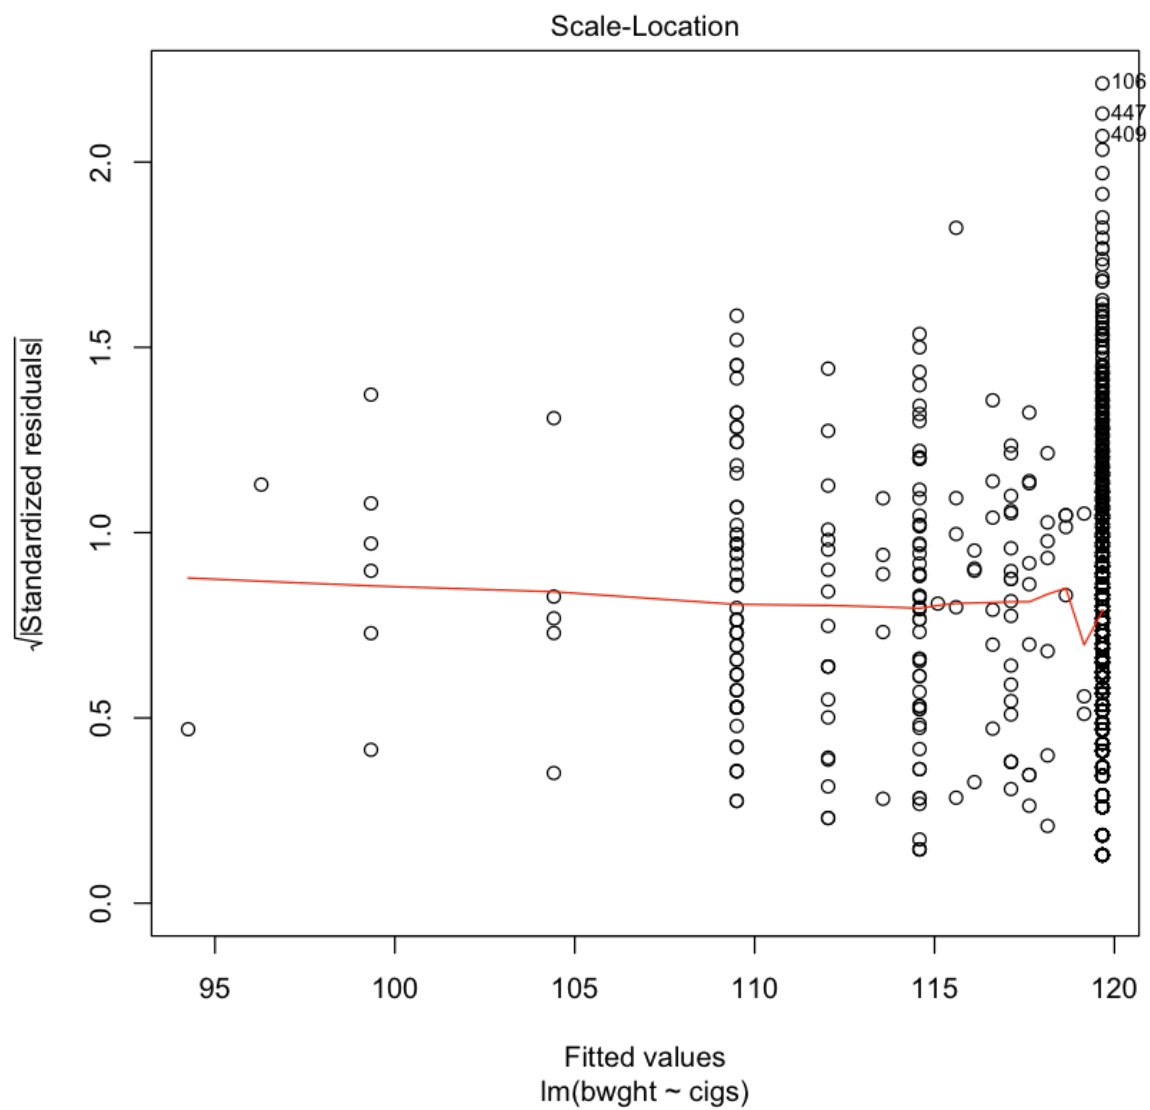
```
d1 <- data[data$bwght>1 & data$bwght<250,]
```

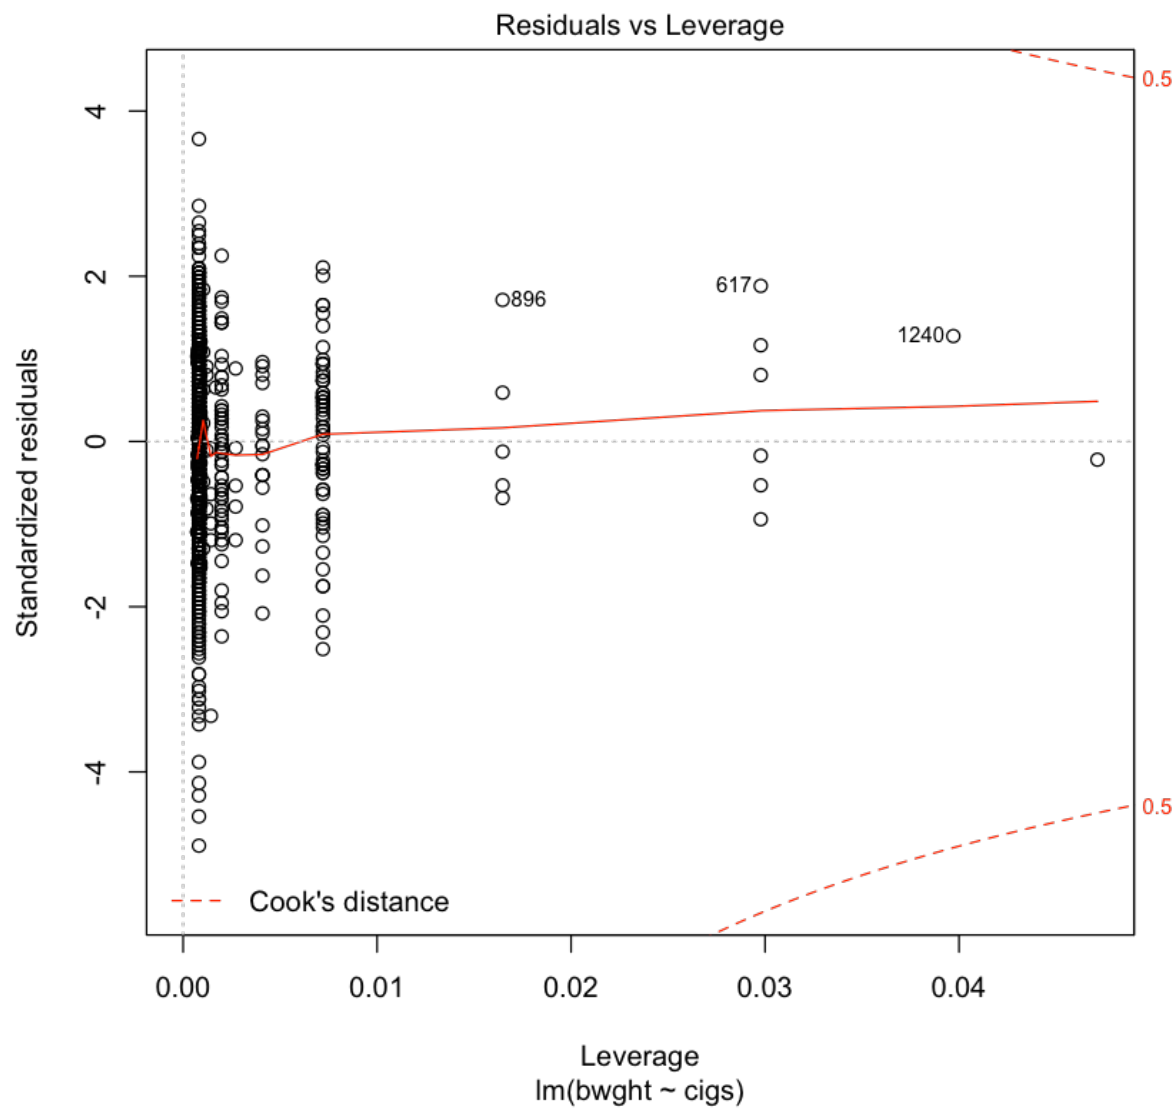
```
m1 <- lm(bwght ~ cigs, data=d1)
```

```
plot(m1)
```









Model for entire sample population

```
summary(m1)
```

Call:

```
lm(formula = bwght ~ cigs, data = d1)
```

Residuals:

Min	1Q	Median	3Q	Max
-96.666	-11.666	0.416	13.334	72.334

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	119.6663	0.5645	211.989	< 2e-16 ***
cigs	-0.5083	0.0889	-5.717	1.32e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.76 on 1375 degrees of freedom

Multiple R-squared: 0.02322, Adjusted R-squared: 0.02251

F-statistic: 32.69 on 1 and 1375 DF, p-value: 1.324e-08

Model for smokers only in sample population

```
m1 <- lm(bwght ~ cigs, data=subset(data, cigs>0))
```

```
summary(m1)
```

Call:

```
lm(formula = bwght ~ cigs, data = subset(data, cigs > 0))
```

Residuals:

Min	1Q	Median	3Q	Max
-62.404	-12.542	0.762	12.010	47.040

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	114.1804	2.4523	46.561	<2e-16 ***
cigs	-0.2220	0.1515	-1.465	0.144

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.13 on 210 degrees of freedom

Multiple R-squared: 0.01012, Adjusted R-squared: 0.005407

F-statistic: 2.147 on 1 and 210 DF, p-value: 0.1443

The model that includes the entire population shows that the intercept is close to the population mean and for every cigarette smoked the birthweight decreases by 0.51 ounces.

The model for smokers out of the sample population indicates a weaker influence among that part of the population with an insignificant p value

The conclusion we can make from this is that the first cigarette is the most meaningful in terms of the differences in birthweight between smokers and non-smokers while pregnant

Question 7:

Now, introduce a new independent variable, faminc, representing family income in thousands of dollars. Examine this variable using the same analysis as in question 3. In addition, produce a scatterplot matrix of bwght, cigs, and faminc. Use the following command (as a starting point):

```
library(car)
scatterplot.matrix(~ bwght + cigs + faminc, data = data2)
```

Note that the car package is needed in order to use the scatterplot.matrix function.

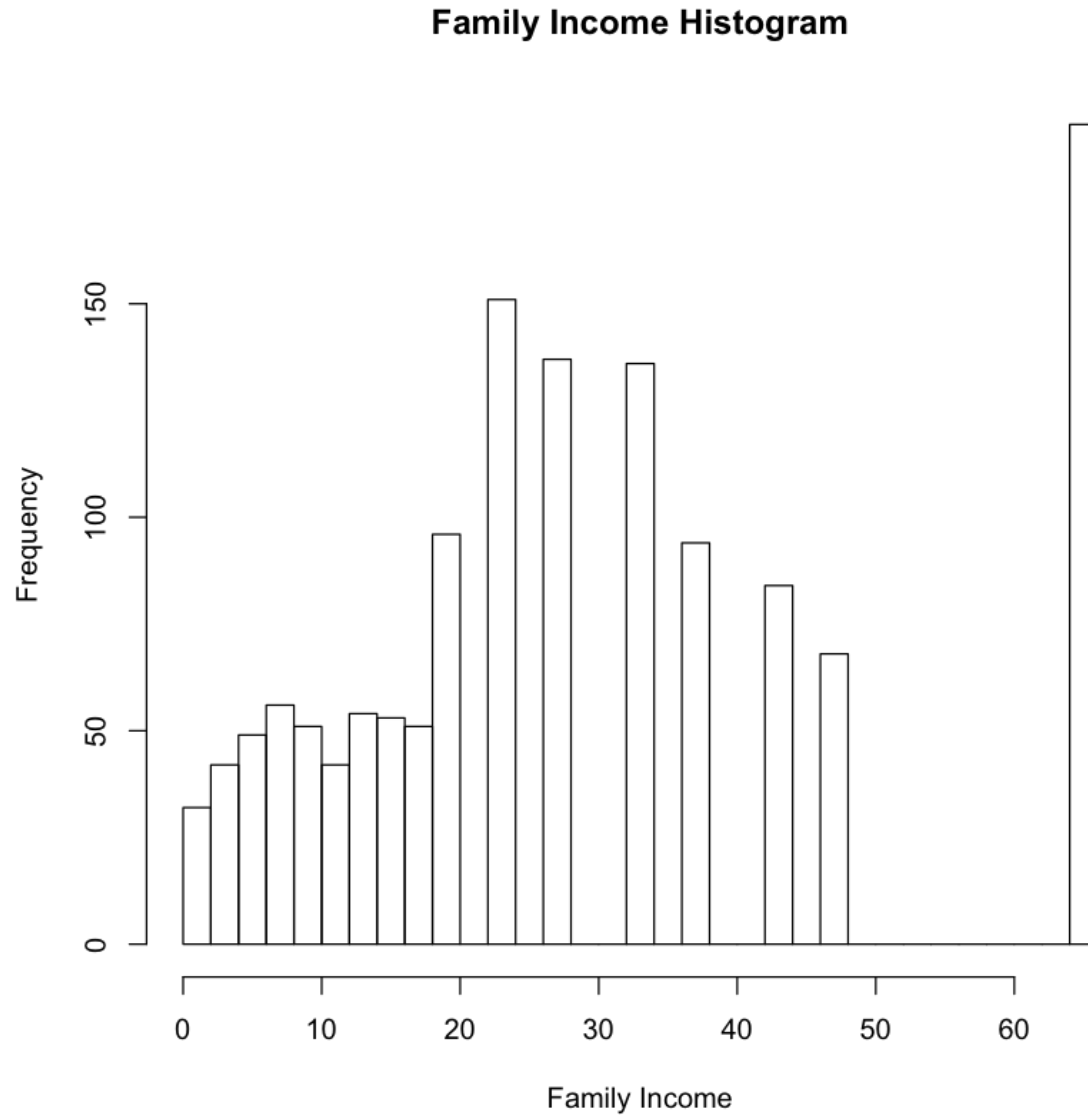
```
summary(data$faminc)
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 0.50    14.50    27.50    29.03   37.50    65.00
```

```
quantile(data$faminc, probs=c(1,5,10,25,50,75,90,99, NA)/100)
```

	quantile	count
1	1%	0.5
2	5%	3.5
3	10%	6.5
4	25%	14.5
5	50%	27.5
6	75%	37.5
7	90%	65
8	99%	65
9	9	NA

```
hist(data$faminc, breaks=30, main="Family Income Histogram",  
xlab="Family Income")
```

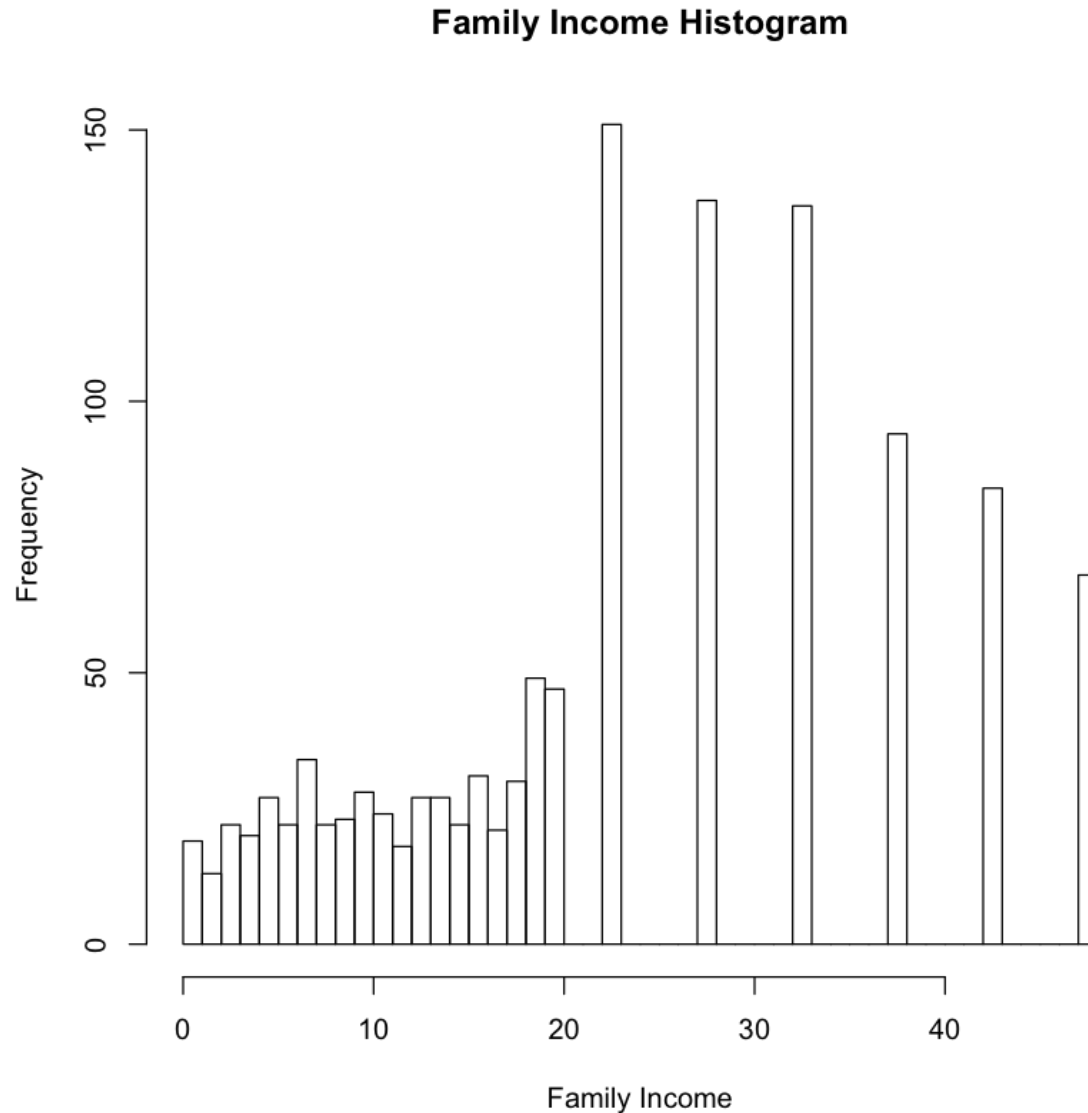



```
table(data$faminc)
```

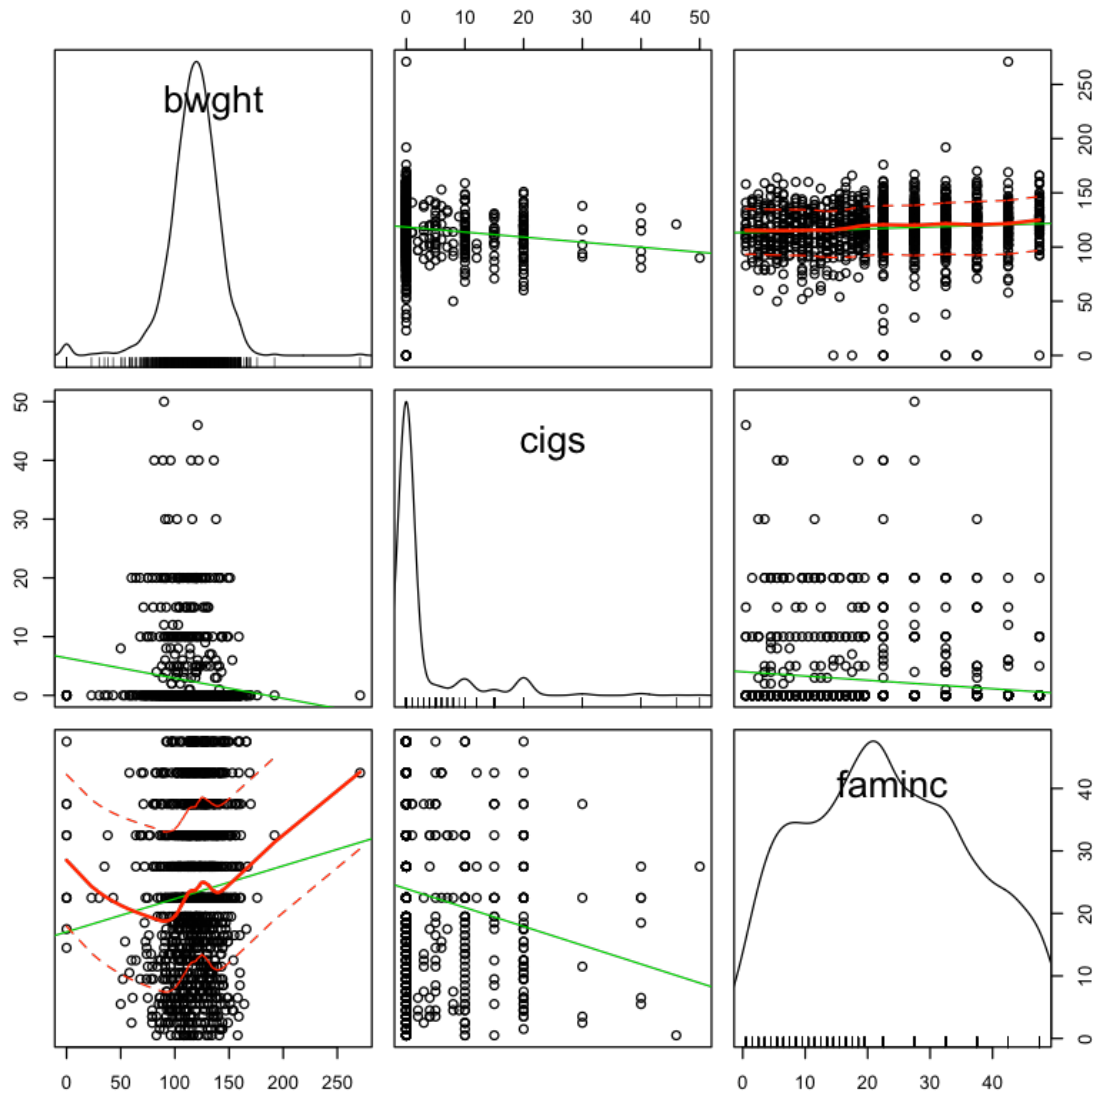
0.5	1.5	2.5	3.5	4.5	5.5	6.5	7.5	8.5	9.5	10.5	11.5	12.5	13.5	14.5	15.5
19	13	22	20	27	22	34	22	23	28	24	18	27	27	22	31
16.5	17.5	18.5	19.5	22.5	27.5	32.5	37.5	42.5	47.5	65					
21	30	49	47	151	137	136	94	84	68	192					

The data seems to indicate a coding error for the value of 65. We see that it is an integer and not a floating point number and there is a large gap between that value and the next lowest. Therefore we will filter out those values greater than 55

```
hist(data$faminc[data$faminc < 55], breaks=50, main="Family Income Histogram", xlab="Family Income")
```



```
scatterplotMatrix(~ bwght + cigs + faminc, data=subset(data,
faminc<55))
```



There appears to be a negative correlation between cigarettes smoked while pregnant and family income

Question 8:

Regress bwght on both cigs and faminc. What coefficient estimates and the standard errors associated with the coefficient estimates do you get? Interpret the results.

```
d2 <- data[data$faminc<55 & data$bwght>1 & data$bwght<250,]
```

```
m2 <- lm(bwght ~ cigs + faminc, data=d2)
```

```
summary(m2)
```

Call:

```
lm(formula = bwght ~ cigs + faminc, data = d2)
```

Residuals:

Min	1Q	Median	3Q	Max
-96.23	-11.35	0.77	13.21	71.21

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	115.70927	1.23968	93.338	< 2e-16	***
cigs	-0.46955	0.09181	-5.114	3.67e-07	***
faminc	0.15646	0.04508	3.471	0.000538	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.8 on 1182 degrees of freedom

Multiple R-squared: 0.03617, Adjusted R-squared: 0.03454

F-statistic: 22.18 on 2 and 1182 DF, p-value: 3.504e-10

Birthweight by cigarette consumption is similar to the first model without faminc, but now we see a positive correlation between faminc and birthweight, $p < .0005$

Question 9:

Explain, in your own words, what the coefficient on cigs in the multiple regression means, and how it is different than the coefficient on cigs in the simple regression? Please provide the intuition to explain the difference, if any.

The cigs coefficient means that, holding the effect of family income constant, each unit of consumption of cigarettes correlates to a drop in birthweight of 0.47 ounces. This is different from the single regression interpretation in that there are no other variable correlations in the model for which to account.

Question 10:

Which coefficient for *cigs* is more negative than the other? Suggest an explanation for why this is so.

The multiple regression model *cigs* coefficient is less negative than the single regression because the multiple regression explains the effects of cigarettes on birthweight after the effect of income has been partialled out